Smart VidiFlow — Master Plan

Canonical as of 2026-09-17. Supersedes the August 20, 2026 Smart VidiFlow Master Plan for architecture locks. Phase status from Aug is retained where still accurate; compute and orchestration layers are brought forward.

Companion lock: `VIDIFLOW-SMART-FLOW-LOCK-2026-09-17.md`
Business pack: `VIDIFLOW-VIDISMART-BUSINESS-PLAN.md`
Revka/Kumiho detail: `REVKA-KUMIHO-CONTEXT-LOCK.md`

Goal

Build Smart VidiFlow: the visual, glance-level control panel for the entire AI agent fleet — agents seen through color, position, and shape, not text — with voice so the operator can talk to any agent and it talks back in real time.

Product surface names: VidiFlow HX (visual agentic OS / Mission Control) + Smart Meeting / Smart App.

Architecture lock (layered — do not merge)

Layer	System	Role
Surface	HX Mission Control + React Flow canvas (+ Peerdom/Rolebase–Maptio DNA)	Visual / HITL only — not the task orchestrator. Agents may appear as board entities.
Orchestration	Hermes Agent Kanban / dispatcher	CORE task orchestration (supersedes Sep 7 “Hermes = entities only”)
Cognition	Revka + Kumiho	Revka = governed agent runtime; Kumiho = graph memory; turn shape engage !’ context !’ answer !’ reflect
Meeting	Lexi + Max	Smart Meeting personas; spine still open (see Open questions)
Data	Postgres + SeaweedFS + AppFlowy + Qdrant	Object store = SeaweedFS (not MinIO). AppFlowy = SOPs / docs / plans / PM. Qdrant = vector knowledge.
Miami compute	Threadripper + 2× Intel Arc Pro B70 + Qwen 3.8 27B	First-class primary Miami inference
Cloud (experimental)	Cloudflare Agent OS (@cloudflare/computer)	Alongside local Hermes — does not replace it

```
flowchart TB
  subgraph Surface["Surface — visual / HITL"]
    MC[HX Mission Control]
    RF[React Flow canvas]
    ENT[Board entities]
  end
  subgraph Orch["Orchestration — CORE"]
    HERMES[Hermes Kanban / dispatcher]
  end
  subgraph Cog["Cognition"]
    REVKA[Revka]
    KUMIHO[Kumiho]
  end
  subgraph Meet["Meeting"]
    LEXI[Lexi]
    MAX[Max]
  end
  subgraph Data["Data"]
    PG[(Postgres)]
    SW[SeaweedFS]
    AF[AppFlowy]
    QD[(Qdrant)]
  end
  subgraph Miami["Miami — primary inference"]
    TR[Threadripper]
    B70[2× Intel B70]
    QW[Qwen 3.8 27B]
  end
  MC --> HERMES
  RF --> HERMES
  ENT -.-> MC
  LEXI & MAX --> HERMES
  HERMES --> REVKA
  REVKA <--> KUMIHO
  REVKA --> QW
  REVKA --> PG & SW & QD
```

```
HERMES --> PG
AF -.-> MC
TR --> B70 --> QW
```

Confirmed decisions

- Name: Smart VidiFlow / VidiFlow HX. “VidiMeet” retired.
- Interface: visual-first; chat in agent drill-down / meeting surfaces, not as the primary panel.
- Task orchestration: Hermes Agent Kanban/dispatcher is core. Mission Control + canvas = visual HITL.
- Agent cognition: Revka (runtime/governance) + Kumiho (graph memory) required for accurate Smart App answers (2026-09-15 lock).
- Local model path (dev / workstation): LM Studio (`http://localhost:1234/v1`), not Ollama. Hermes context floor "64K where Hermes profiles are used.
- Miami primary brain: Qwen 3.8 27B on Threadripper + 2× Intel Arc Pro B70.
- Voice / meetings (interim): LiveKit self-hosted + Agents + Piper / Edge TTS still in the stack; Lexi/Max UI in progress. Arch commit still open: Tavus CVI vs LiveKit + sovereign STT.
- Knowledge: Qdrant + FastEmbed; per-tenant namespaces `kb_{tenant_id}`.
- Object storage: SeaweedFS S3-compatible (e.g. :9000). Do not pull MinIO (closed source).
- Cloud media: `cdn.vidi.news` = Cloudflare R2 bucket `vidismart`.
- Docs / PM: AppFlowy (not Notion) for SOPs, docs, plans, project management.
- Cloud agents: `@cloudflare/computer` experimental track beside local Hermes fleet.
- Foundation reused: `M:\code\vidismart\agent-backend` (:3007), agent UI, piper, `.env.tts`.

Implementation status (carried forward from Aug 2026)

Phase 1 — Hermes + local models — IN PROGRESS

- Hermes installed/verified historically; profiles = one per agent role where Hermes is the worker surface.
- Bring-forward: Hermes Kanban/dispatcher is the orchestration layer; pair workers with Revka+Kumiho for governed accurate turns.
- LM Studio endpoint (not Ollama) for local OpenAI-compatible inference; Miami B70/Qwen is the production Miami target.

Phase 2 — Control plane (agent-backend) — MOSTLY DONE

- Fleet registry, telemetry, command dispatch, Socket.IO events.
- Register Hermes (and later cloud) agents through fleet API so Mission Control shows one grid.

Phase 3 — Visual control panel — DONE (evolve to HX Mission Board)

- agent-ui / HX Mission Control + React Flow canvas path.
- Keep visual/HITL; do not make the canvas the task DB — that is Hermes Kanban.

Phase 4 — Voice / Smart Meeting — SERVER PARTIAL, AGENTS PENDING

- LiveKit server path exists; Lexi + Max meeting layer started 2026-09-16.
- Open: Tavus CVI vs LiveKit face-only + BYO/sovereign STT.

Phase 5 — Knowledge — DONE (custom pipeline)

- Qdrant + FastEmbed ingest; wire citations into Revka/Kumiho-aware agent prompts.

Phase 6 — Storage — DONE (SeaweedFS)

- SeaweedFS all-in-one / S3 path. MinIO removed from current plan.

Phase 7 — Cloudflare Agent OS — EXPERIMENTAL

- Durable Object agents; feature-flagged; local Hermes remains production orchestration path.

Phase 8 — Revka + Kumiho (NEW / REQUIRED)

- Wire engage !ƈontext !ąnswer !ɾeflect into final Smart App.
- Glance UI must render live Kumiho graph state, not mock nodes.
- Detail: `REVKA-KUMIHO-CONTEXT-LOCK.md`.

Phase 9 — Miami fleet (NEW / LOCKED)

- Stand up / document Threadripper + 2× Intel Arc Pro B70 serving Qwen 3.8 27B.
- Treat as first-class in bandwidth / privacy / demo story (separate from AirPMD dual RTX PRO 6000 lean MVP track).

Compute plan (updated)

Miami — PRIMARY (locked 2026-09-17)

- Hardware: Threadripper + 2× Intel Arc Pro B70
- Model: Qwen 3.8 27B
- First-class in smart-flow structure and sovereign demos.

Workstation / lab — retained

- Local LM Studio brains on available NVIDIA cards (e.g. RTX 4090 path from Aug plan) remain valid for dev/lab.
- Planned 2× RTX 6000 Pro expansion (AirPMD / high-VRAM track) stays a separate scale ladder — does not replace Miami B70/Qwen as the Miami lock.

Cloud inference — optional

- Workers AI / Paid plan models for Cloudflare Agent OS only; not a substitute for Miami sovereign path.

Validation checklist (updated)

- [ ] Hermes Kanban/dispatcher owns task orchestration; Mission Control is visual/HITL only
- [ ] Revka + Kumiho engage!'reflectgreen on Smart App agents
- [ ] SeaweedFS serving media/object buckets (no MinIO)
- [ ] Miami Threadripper + 2× B70 serving Qwen 3.8 27B end-to-end
- [ ] Fleet dashboard shows real Hermes (+ optional cloud) agents
- [ ] Lexi/Max meeting path works for voice (arch TBD)
- [ ] Document ingest !Qdrant searchable with tenant isolation
- [ ] AppFlowy holds SOPs/plans as team knowledge base

Risks

- Layer collapse — merging Hermes, Mission Control, and Revka into one blob. Keep layers separate.

- Ollama drift — LM Studio (or Miami OpenAI-compatible endpoint) is the local path; do not reintroduce Ollama as default.
- MinIO reintroduction — closed source; SeaweedFS is the lock.
- 64K context floor — set explicitly wherever Hermes + LM Studio are used.
- Cloudflare computer preview — feature-flag only.
- Meeting arch indecision — Tavus vs LiveKit+sovereign STT blocks a final Smart Meeting commit.

Open questions

1. Meeting spine: Tavus CVI vs LiveKit + sovereign STT — pick for production.
2. Org map DNA: Peerdom license vs Maptio/Rolebase OSS fork for production map sync.
3. Full-duplex vs push-to-talk for first operator voice sessions.
4. Multi-operator / role-based access to Mission Control.

Document history

Date	Change
2026-08-20	Original Smart VidiFlow Master Plan (Hermes workers, LM Studio, SeaweedFS, LiveKit, Cloudflare experimental, 4090 / RTX 6000 tiers)
2026-09-15	Revka + Kumiho required for Smart App accuracy
2026-09-17	Hermes = core task orchestration; Mission Control = visual/HITL; SeaweedFS confirmed vs MinIO; Miami TR + 2× B70 + Qwen 3.8 27B; AppFlowy KB; this master plan brought forward

VidiSmart — Confidential Internal Document

Smart VidiFlow Master Plan | Updated 2026-09-17